

Lab 1: Stock Price Prediction Analytics using Snowflake & Airflow

Abhinita Sanabada
(018320874)

Rishitha Gogineni
(019162377)

October 5, 2025

1 Team Introduction

- **Abhinita Sanabada** (018320874): End-to-end pipeline owner. Implemented Airflow DAGs, Snowflake SQL/transactions, and report authoring.
- **Rishitha Gogineni** (019162377): Implemented Sql queries in snowflake with understanding table structure in yfinance.

Abstract

We implemented a secure, reproducible stock analytics workflow using Airflow, yfinance, and Snowflake. An ETL DAG ingests OHLCV data into `RAW.STOCK_PRICES`. A second DAG trains a Snowflake-native `SNOWFLAKE.ML.FORECAST` model and writes predictions to `MODEL.FORECASTS`. A final table, `ANALYTICS.FINAL_PRICES_FORECAST`, unions actuals and forecasts for downstream visualization. All Snowflake credentials (account, user, password, role, warehouse, database) are stored only in Airflow Connections, and pipeline parameters are managed via Airflow Variables.

2 Problem Statement

Build an end-to-end analytics pipeline that:

1. Extracts daily OHLCV for selected tickers via yfinance.
2. Forecasts daily close prices using Snowflake's built-in ML forecasting.
3. Unifies actuals and forecasts in a single analytics table.
4. Uses Airflow for orchestration, SQL/Python transactions for correctness, and Airflow Connections/Variables for secure configuration.
5. Produces reproducible runs, screenshots, and a public code repository.

Success criteria: both DAGs succeed; RAW, MODEL, ANALYTICS populated; final table supports plotting Actual vs. Forecast; screenshots and repo links are provided.

```

8      -- RAW Layer: Ingested daily bars
9      CREATE TABLE IF NOT EXISTS RAW.STOCK_PRICES (
10         SYMBOL          STRING          NOT NULL,
11         TS              TIMESTAMP_NTZ NOT NULL,
12         OPEN            FLOAT,
13         HIGH            FLOAT,
14         LOW             FLOAT,
15         CLOSE           FLOAT,
16         ADJ_CLOSE       FLOAT,
17         VOLUME          NUMBER(38,0),
18         LOAD_TS         TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),
19         CONSTRAINT PK_STOCK_PRICES PRIMARY KEY (SYMBOL, TS)
20     );
21
22     -- MODEL Layer: Forecast outputs (one row per date per symbol)
23     CREATE TABLE IF NOT EXISTS MODEL.FORECASTS (
24         SYMBOL          STRING          NOT NULL,
25         TS              DATE            NOT NULL,
26         PREDICTED_CLOSE FLOAT          NOT NULL,
27         MODEL_NAME      STRING          NOT NULL,
28         TRAINED_AT      TIMESTAMP_NTZ NOT NULL,
29         HORIZON_D        NUMBER(5,0)   NOT NULL,
30         LOAD_TS         TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),
31         CONSTRAINT PK_FORECASTS PRIMARY KEY (SYMBOL, TS, MODEL_NAME)
32     );
33
34     -- ANALYTICS Layer: Union of actuals + forecasts (latest partition wins on overlap)
35     CREATE TABLE IF NOT EXISTS ANALYTICS.FINAL_PRICES_FORECAST (
36         SYMBOL          STRING          NOT NULL,
37         TS              DATE            NOT NULL,
38         CLOSE           FLOAT,
39         SOURCE          STRING          NOT NULL, -- 'ACTUAL' or 'FORECAST'
40         MODEL_NAME      STRING,         -- null for ACTUAL
41         LOAD_TS         TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),

```

Figure 1: snowflake table create

3 Dataset(s)

Source: `yfinance` Python library; daily OHLCV pulls via Airflow DAG #1.

Universe: Symbols configured via Airflow Variable `stock_symbols`. Example used: AAPL, MSFT, TSLA.

Time Window: `lookback_days` ending on DAG run date (default: 365).

Granularity / Fields: Daily bars with columns OPEN, HIGH, LOW, CLOSE, ADJ_CLOSE, VOLUME.

Quality & Preprocessing: Idempotent loads with MERGE on composite key (SYMBOL, TS); timestamps normalized to Snowflake `TIMESTAMP_NTZ`; missing market days left as gaps (no forward fill).

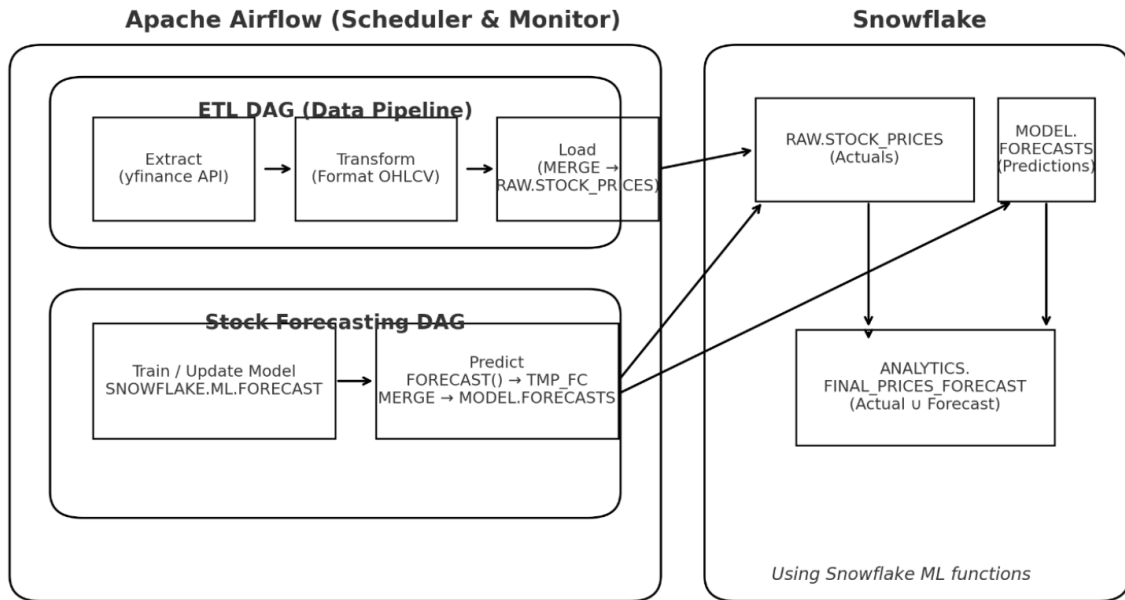


Figure 2: Architecture Diagram

4 System Architecture

Overview

We use three schemas inside `USER_DB_CATFISH`: **RAW** (ingest), **MODEL** (predictions), **ANALYTICS** (consumption). Two DAGs orchestrate the flow:

- **DAG #1 `yfinance_etl`**: downloads OHLCV for `stock_symbols`, MERGEs into `RAW.STOCK_PRICES`
- **DAG #2 `ml_forecast`**: trains/updates `SNOWFLAKE.ML.FORECAST` on multi-series history; writes predictions to `MODEL.FORECASTS`; *unions* actuals + forecasts into `ANALYTICS.FINAL_PRICES_FORECAST`

Architecture Diagram

Screenshots (to be included)

- **Figure 2**: Airflow DAGs list showing both DAGs present.
- **Figure 3**: `yfinance_etl` Grid/Graph view with successful run.
- **Figure 4**: `ml_forecast` Grid/Graph view with successful run.

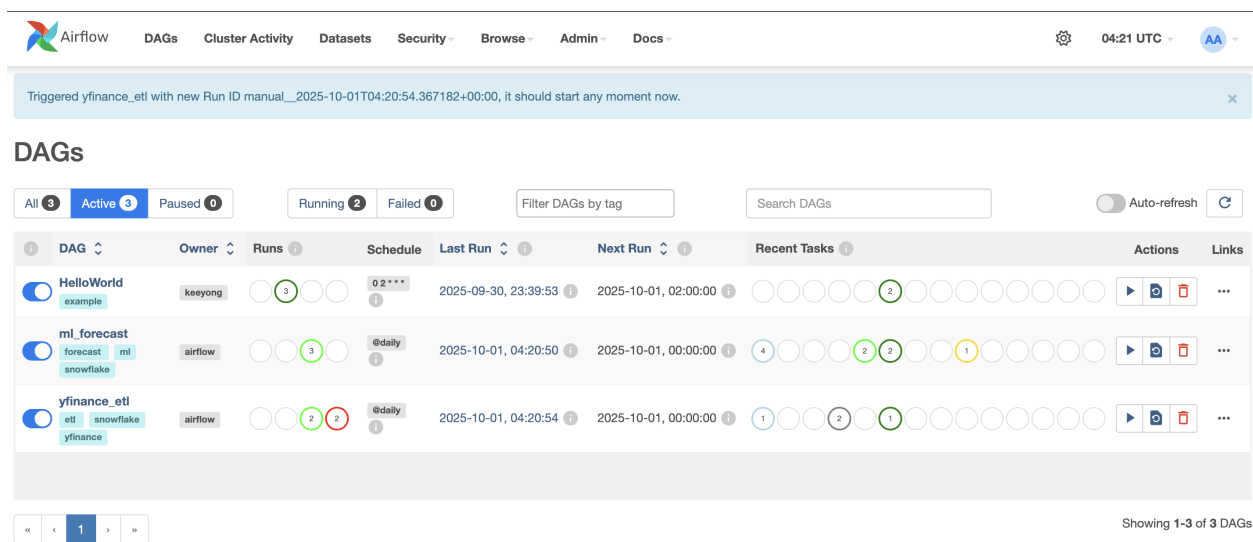


Figure 3: Airflow dags screen shot

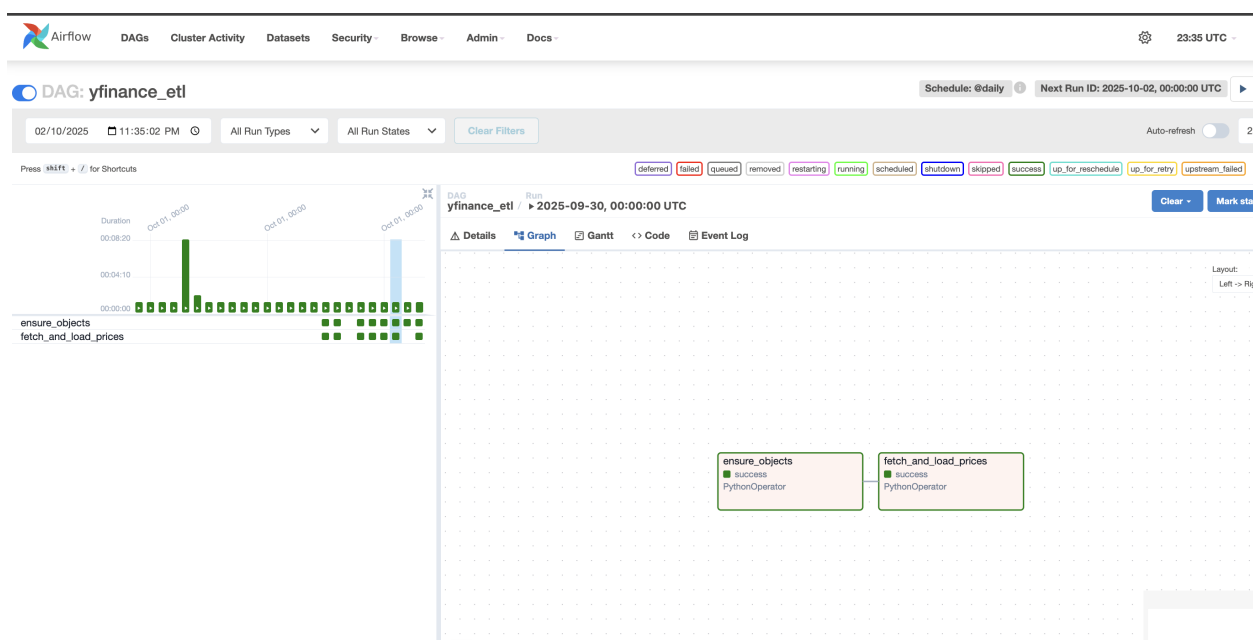


Figure 4: yfinance etl Grid/Graph

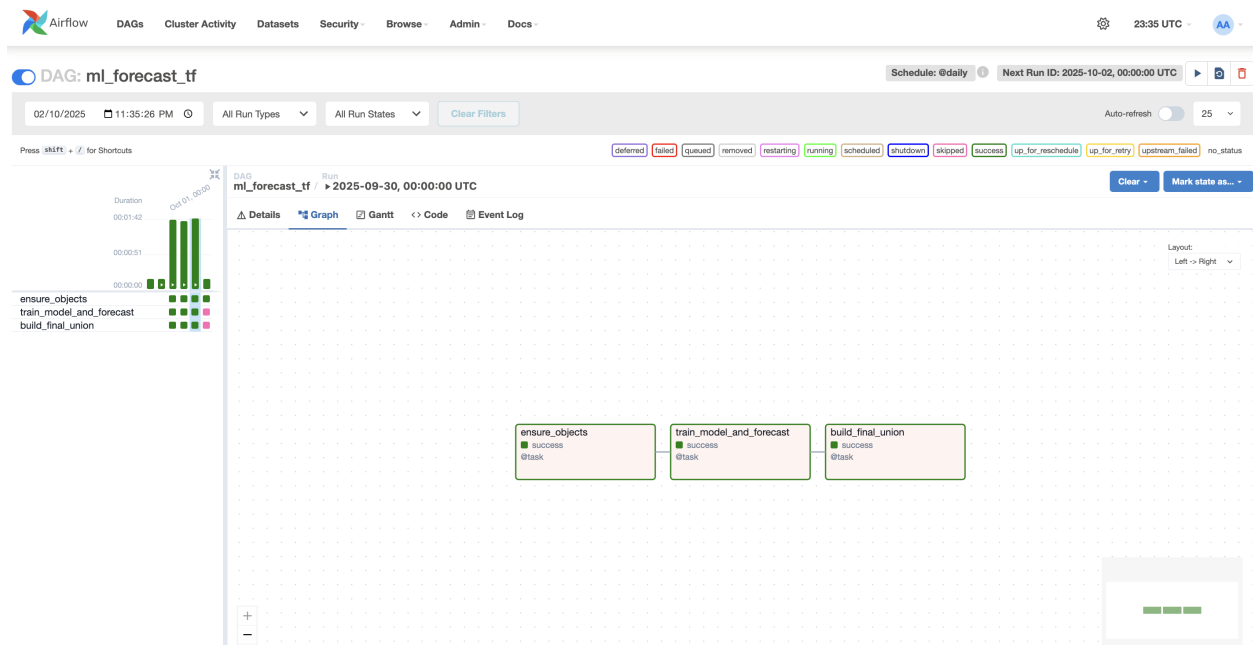


Figure 5: ml forecast Grid/Graph

5 Tables

5.1 Sample analytics output (Actual vs. Forecast)

Table 1: Sample rows from `ANALYTICS.FINAL_PRICES_FORECAST`. Replace with real rows.

SYMBOL	TS	CLOSE	SOURCE	MODEL_NAME
AAPL	2025-09-15	228.34	ACTUAL	—
AAPL	2025-10-05	231.10	FORECAST	SNOWFLAKE_ML
MSFT	2025-09-15	422.90	ACTUAL	—
MSFT	2025-10-05	427.50	FORECAST	SNOWFLAKE_ML

5.2 Schema summary (brief)

Table 2: Key columns in `RAW.STOCK_PRICES`.

Column	Type	Notes
SYMBOL	STRING (PK)	Ticker symbol
TS	TIMESTAMP_NTZ (PK)	Bar timestamp
OPEN,HIGH,LOW,CLOSE	FLOAT	Daily OHLC
ADJ_CLOSE	FLOAT	Adjusted close
VOLUME	NUMBER(38,0)	Shares traded

6 Data Model

All objects live in database USER_DB_CATFISH (warehouse: CATFISH_QUERY_WH, both configured via Airflow Connection). Schemas: RAW, MODEL, ANALYTICS.

RAW.STOCK_PRICES

```
SYMBOL STRING NOT NULL; TS TIMESTAMP_NTZ NOT NULL; OPEN FLOAT; HIGH  
FLOAT; LOW FLOAT; CLOSE FLOAT; ADJ_CLOSE FLOAT; VOLUME NUMBER(38,0);  
LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP();
```

Primary key: (SYMBOL, TS).

MODEL.FORECASTS

```
SYMBOL STRING NOT NULL; TS DATE NOT NULL; PREDICTED_CLOSE FLOAT NOT  
NULL; MODEL_NAME STRING NOT NULL; TRAINED_AT TIMESTAMP_NTZ NOT NULL;  
HORIZON_D NUMBER(5,0) NOT NULL; LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP();
```

Primary key: (SYMBOL, TS, MODEL_NAME).

ANALYTICS.FINAL_PRICES_FORECAST

```
SYMBOL STRING NOT NULL; TS DATE NOT NULL; CLOSE FLOAT; SOURCE STRING  
NOT NULL; MODEL_NAME STRING; LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP();
```

Primary key: (SYMBOL, TS, SOURCE); SOURCE ∈ {ACTUAL, FORECAST}.

One-time Bootstrap DDL

Listing 1: Bootstrap DDL (run once).

```
1 BEGIN;  
2 CREATE SCHEMA IF NOT EXISTS RAW;  
3 CREATE SCHEMA IF NOT EXISTS MODEL;  
4 CREATE SCHEMA IF NOT EXISTS ANALYTICS;  
5  
6 CREATE TABLE IF NOT EXISTS RAW.STOCK_PRICES (  
7     SYMBOL STRING NOT NULL, TS TIMESTAMP_NTZ NOT NULL,  
8     OPEN FLOAT, HIGH FLOAT, LOW FLOAT, CLOSE FLOAT, ADJ_CLOSE FLOAT,  
9     VOLUME NUMBER(38,0),  
10    LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),  
11    CONSTRAINT PK_STOCK_PRICES PRIMARY KEY (SYMBOL, TS)  
12 );  
13  
14 CREATE TABLE IF NOT EXISTS MODEL.FORECASTS (  
15     SYMBOL STRING NOT NULL, TS DATE NOT NULL, PREDICTED_CLOSE FLOAT  
16     NOT NULL,
```

```

15  MODEL_NAME STRING NOT NULL, TRAINED_AT TIMESTAMP_NTZ NOT NULL,
    HORIZON_D NUMBER(5,0) NOT NULL,
16  LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP(),
17  CONSTRAINT PK_FORECASTS PRIMARY KEY (SYMBOL, TS, MODEL_NAME)
18 );
19
20 CREATE TABLE IF NOT EXISTS ANALYTICS.FINAL_PRICES_FORECAST (
21  SYMBOL STRING NOT NULL, TS DATE NOT NULL, CLOSE FLOAT,
22  SOURCE STRING NOT NULL, MODEL_NAME STRING, LOAD_TS TIMESTAMP_NTZ
    DEFAULT CURRENT_TIMESTAMP(),
23  CONSTRAINT PK_FINAL PRIMARY KEY (SYMBOL, TS, SOURCE)
24 );
25 COMMIT;

```

7 Implementation

7.1 Airflow Connections & Variables

We created a Snowflake Connection **snowflake_catfish** with account, user, password, role, default warehouse CATFISH_QUERY_WH, and database USER_DB_CATFISH. **No secrets in code.** Variables:

- `stock_symbols`: JSON list, e.g., `["AAPL", "MSFT", "TSLA"]`.
- `lookback_days`: e.g., 365.
- `forecast_horizon_days`: e.g., 14.
- `target_schema_raw`=RAW, `target_schema_model`=MODEL, `target_schema_analytics`=ANALYTICS

7.2 DAG #1: yfinance_etl (ETL)

Python downloads OHLCV for `stock_symbols` over the last `lookback_days` and MERGES into `RAW.STOCK_PRICES`. Credentials/DB/WH/role come from the Airflow Connection at run-time (via `snowflake.connector`). We use a transactional pattern (commit/rollback).

7.3 Create Snowflake Connection in Airflow UI

Step 1: Open the Airflow web UI (e.g., `http://localhost:8080`).

Step 2: Navigate to **Admin** → **Connections**.

Step 3: Click + (*Add a new record*).

Step 4: Fill the form:

- **Conn Id:** `snowflake_catfish`
- **Conn Type:** Snowflake

The screenshot shows the Airflow web interface with the 'Edit Connection' form. The form fields are as follows:

Field	Value
Connection Id *	snowflake_catfish
Connection Type *	Snowflake
Description	
Schema	RAW
Login	CATFISH
Password	snowflake password

A warning message at the top states: "Warning: Fields that are currently populated can be modified but cannot be deleted. To delete data from a field, delete the Connection object and create a new one."

Figure 6: Airflow Connection

- **Login:** your Snowflake *username*
- **Password:** your Snowflake *password*
- **Extra (JSON):** paste the following and adjust values:

```
{
  "account": "abcd-xy123",
  "warehouse": "COMPUTE_WH",
  "database": "YOUR_DB",
  "schema": "RAW",
  "role": "SYSADMIN"
}
```

Step 5: Click **Test** (top-right). If successful, click **Save**.

Notes.

- The DAGs access this via `BaseHook.get_connection("snowflake_catfish")` and read fields from `c.extra_dejson`.
- If your ML pipeline expects a different default schema (e.g., `MODEL`), either change it here or pass a schema argument in code.

7.4 Create Airflow Variables

Step 1: Go to **Admin** → **Variables**.

Step 2: Click **+** to add each key/value below (use exact keys):

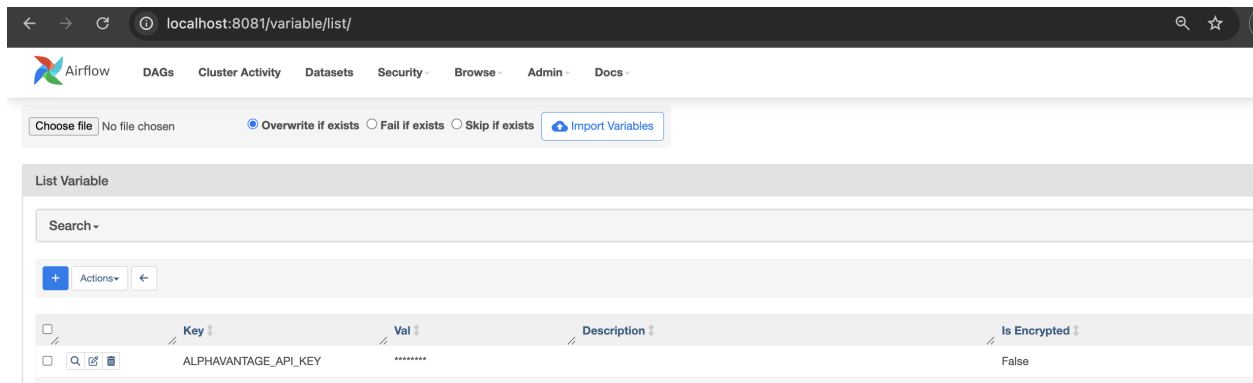


Figure 7: Airflow variable

Runtime DML (executed by the DAG after staging rows):

SQL

Listing 2: ETL MERGE (executed each run).

```

1
2 BEGIN;
3 CREATE TEMP TABLE TMP_LOAD (
4     SYMBOL STRING, TS TIMESTAMP_NTZ, OPEN FLOAT, HIGH FLOAT, LOW FLOAT
5     ,
6     CLOSE FLOAT, ADJ_CLOSE FLOAT, VOLUME NUMBER(38,0)
7 );
8 -- Python inserts many rows into TMP_LOAD via executemany(...)
9 MERGE INTO RAW.STOCK_PRICES AS t
10 USING TMP_LOAD AS s
11     ON t.SYMBOL = s.SYMBOL
12     AND t.TS = s.TS
13 WHEN MATCHED THEN UPDATE SET
14     OPEN=s.OPEN, HIGH=s.HIGH, LOW=s.LOW, CLOSE=s.CLOSE,
15     ADJ_CLOSE=s.ADJ_CLOSE, VOLUME=s.VOLUME, LOAD_TS=CURRENT_TIMESTAMP
16     ()
17 WHEN NOT MATCHED THEN INSERT (
18     SYMBOL, TS, OPEN, HIGH, LOW, CLOSE, ADJ_CLOSE, VOLUME
19 ) VALUES (
20     s.SYMBOL, s.TS, s.OPEN, s.HIGH, s.LOW, s.CLOSE, s.ADJ_CLOSE, s.
21     VOLUME
22 );
23 COMMIT;

```

7.5 DAG #2: ml_forecast (Snowflake ML) + Final Union

This DAG uses only `SnowflakeOperator`. It trains/updates a multi-series model via `SNOWFLAKE.ML.FORECAST`, stages horizon-wide predictions, MERGEs them into `MODEL.FORECASTS`, and rebuilds `ANALYTICS.FINAL` by unioning `ACTUAL` and `FORECAST` rows.

Model Training, Forecasting & Upsert

Listing 3: Snowflake ML model + forecast + upsert.

```
1 BEGIN;
2 USE SCHEMA MODEL;
3
4 WITH symbols AS (
5     SELECT value::string AS symbol
6     FROM TABLE(FLATTEN(input => PARSE_JSON('{{ var.value.stock_symbols
7         }}'))))
8
9 training_data AS (
10     SELECT
11         TO_VARIANT(sp.SYMBOL) AS SERIES,
12         sp.TS,
13         sp.CLOSE
14     FROM RAW.STOCK_PRICES sp
15     JOIN symbols s ON s.symbol = sp.SYMBOL
16     WHERE sp.TS >= DATEADD('day', -{{ var.value.lookback_days |
17         default('365', true) }}, CURRENT_TIMESTAMP())
18 )
19
20 CREATE OR REPLACE SNOWFLAKE.ML.FORECAST PRICE_FORECASTER (
21     INPUT_DATA => SYSTEM$QUERY_REFERENCE($$ SELECT SERIES, TS,
22         CLOSE FROM training_data $$),
23     SERIES_COLNAME => 'SERIES',
24     TIMESTAMP_COLNAME => 'TS',
25     TARGET_COLNAME => 'CLOSE',
26     CONFIG_OBJECT => {{ '{{' }} 'method':'fast','on_error':'skip'
27         '{{ '}}' }}
28 );
29
30 CREATE OR REPLACE TEMP TABLE TMP_FC AS
31 SELECT
32     SERIES::STRING AS SYMBOL,
33     CAST(TS AS DATE) AS TS,
34     FORECAST AS PREDICTED_CLOSE,
35     'SNOWFLAKE_ML' AS MODEL_NAME,
```

```

32     CURRENT_TIMESTAMP ()                AS TRAINED_AT,
33     {{ var.value.forecast_horizon_days | default('14', true) }}::
        NUMBER AS HORIZON_D
34 FROM TABLE(PRICE_FORECASTER!FORECAST(
35     FORECASTING_PERIODS => {{ var.value.forecast_horizon_days |
        default('14', true) }}
36 ));
37
38 MERGE INTO MODEL.FORECASTS AS t
39 USING TMP_FC AS s
40     ON t.SYMBOL      = s.SYMBOL
41     AND t.TS          = s.TS
42     AND t.MODEL_NAME = s.MODEL_NAME
43 WHEN MATCHED THEN UPDATE SET
44     PREDICTED_CLOSE = s.PREDICTED_CLOSE,
45     TRAINED_AT      = s.TRAINED_AT,
46     HORIZON_D       = s.HORIZON_D,
47     LOAD_TS         = CURRENT_TIMESTAMP ()
48 WHEN NOT MATCHED THEN INSERT (
49     SYMBOL, TS, PREDICTED_CLOSE, MODEL_NAME, TRAINED_AT, HORIZON_D
50 ) VALUES (
51     s.SYMBOL, s.TS, s.PREDICTED_CLOSE, s.MODEL_NAME, s.TRAINED_AT, s.
        HORIZON_D
52 );
53 COMMIT;

```

Final Union Build (ACTUAL FORECAST)

Listing 4: Rebuild ANALYTICS final table.

```

1 BEGIN;
2 USE SCHEMA ANALYTICS;
3
4 WITH symbols AS (
5     SELECT value::string AS symbol
6     FROM TABLE(FLATTEN(input => PARSE_JSON('{{ var.value.stock_symbols
        }}'))))
7 )
8
9 TRUNCATE TABLE ANALYTICS.FINAL_PRICES_FORECAST;
10
11 -- ACTUALS from RAW
12 INSERT INTO ANALYTICS.FINAL_PRICES_FORECAST (SYMBOL, TS, CLOSE,
        SOURCE, MODEL_NAME)

```

```

13 SELECT
14     sp.SYMBOL,
15     CAST(sp.TS AS DATE) AS TS,
16     sp.CLOSE,
17     'ACTUAL' AS SOURCE,
18     NULL AS MODEL_NAME
19 FROM RAW.STOCK_PRICES sp
20 JOIN symbols s ON s.symbol = sp.SYMBOL;
21
22 -- FORECASTS from MODEL
23 INSERT INTO ANALYTICS.FINAL_PRICES_FORECAST (SYMBOL, TS, CLOSE,
24     SOURCE, MODEL_NAME)
25 SELECT
26     f.SYMBOL,
27     f.TS,
28     f.PREDICTED_CLOSE AS CLOSE,
29     'FORECAST' AS SOURCE,
30     f.MODEL_NAME
31 FROM MODEL.FORECASTS f
32 JOIN symbols s ON s.symbol = f.SYMBOL;
33 COMMIT;

```

7.6 Transactions & Error Handling

- **DAG #1:** `snowflake.connector` uses `conn.autocommit(False)` and wraps DDL/DML in `try/except` with `commit()` on success and `rollback()` on failure.
- **DAG #2:** Each `SnowflakeOperator` task uses `BEGIN; ...COMMIT;` so the entire step is atomic.

8 Results

After triggering `yfinance_etl` and then `ml_forecast`, we validated:

- `RAW.STOCK_PRICES` contains daily OHLCV for each ticker.
- `MODEL.FORECASTS` contains `forecast_horizon_days` predictions per ticker with `MODEL_NAME='SNOWFLAKE_ML'`.
- `ANALYTICS.FINAL_PRICES_FORECAST` holds both ACTUAL and FORECAST rows.

Validation Queries

```

1 SELECT 'RAW' AS t, COUNT(*) c FROM RAW.STOCK_PRICES
2 UNION ALL SELECT 'MODEL', COUNT(*) FROM MODEL.FORECASTS
3 UNION ALL SELECT 'ANALYTICS', COUNT(*) FROM ANALYTICS.
   FINAL_PRICES_FORECAST;
4
5 SELECT SYMBOL, MIN(TS) AS min_ts, MAX(TS) AS max_ts, COUNT(*) AS n
6 FROM ANALYTICS.FINAL_PRICES_FORECAST
7 GROUP BY SYMBOL
8 ORDER BY SYMBOL, min_ts;

```

9 Discussion

Pros. Snowflake-side training avoids heavy Python deps and simplifies daily orchestration. Airflow manages schedule, retries, and secret handling via Connections/Variables. Transactions ensure atomic loads and clear failure modes.

create the Airflow Connection and Variables

SYMBOL STRING NOT NULL; TS DATE NOT NULL; CLOSE FLOAT; SOURCE STRING NOT NULL; MODEL_NAME STRING; LOAD_TS TIMESTAMP_NTZ DEFAULT CURRENT_TIMESTAMP (

Primary key: (SYMBOL, TS, SOURCE); SOURCE \in {ACTUAL, FORECAST}.

Airflow Setup: Snowflake Connection & Variables

Required

- `stock_symbols` (JSON):
["AAPL", "MSFT", "TSLA"]
- `lookback_days`: 365
- `target_schema_raw`: RAW

Optional (recommended)

- `forecast_horizon_days`: 14
- `target_schema_model`: MODEL
- `target_schema_analytics`: ANALYTICS

Tips.

- Ensure JSON variables are valid (double quotes, no trailing commas).
- Access in code with `Variable.get("stock_symbols");` parse JSON if needed.
- UI changes to Connections/Variables take effect immediately; new DAG code may need a short scheduler refresh window.

10 Future Work

- **Backtesting & metrics:** Rolling-origin evaluation (MAPE, RMSE) and a performance table per symbol.
- **Richer features:** Holiday calendars, volatility features, and exogenous regressors (e.g., sector ETFs, macro rates).
- **Ops:** SLA alerts on task duration, failure notifications, CI linting for DAGs, and unit tests for SQL.
- **Data quality:** Great Expectations/dbt tests for freshness, nulls, and monotonic timestamps.
- **Productization:** Daily schedules with retries and timeouts; dashboard for run health and data latency.

11 Quick Checklist

- `snowflake_catfish` exists and **Test** passes.
- `stock_symbols`, `lookback_days`, and `target_schema_raw` are present.
- DAG code references the same *Conn Id* and variable names.

12 Common Issues & Fixes

- **“The conn_id isn’t defined”:** the Connection ID in UI doesn’t match the DAG.
- **Auth failures:** wrong account/role/warehouse in Extra.
- **JSON errors in Variables:** re-save with valid JSON (use the code blocks above).

13 Conclusion

The pipeline meets the lab goals: ETL → Snowflake ML Forecast → union table; clean Airflow orchestration with transactions; secure configuration using Connections/Variables.

References

- yfinance (PyPI): <https://pypi.org/project/yfinance/>
- Apache Airflow: <https://airflow.apache.org/>
- Snowflake Documentation: <https://docs.snowflake.com/>